

Kaggle Competition –

Dear Students,

With this challenge, we invite you to dive into the **Company Insolvency Dataset**, explore the available data, and build a predictive model to forecast company insolvency as accurately as possible.

Your Tasks:

1. **Build a working model** that can predict company insolvency.
2. **(Optional but encouraged):** Deploy your model locally as an API with a streamlit UI.
3. **(Required):** Deploy your final model to **Render** as you did it with the web dashboard.

Guidelines:

- Use **Python** for this project.
- Apply **EDA techniques** to understand the dataset and guide your modeling.
- Include **feature selection** and methods to **avoid overfitting**.
- You are free to use any Python libraries (e.g., **PyCaret**) during development.
- The **final deployed model must be implemented using sklearn or another Python SDK-based library** so it can be programmatically called.
- Leverage everything you've learned over the past weeks — from data preprocessing to model tuning and deployment.

Objective:

This competition is not just about accuracy — it's about demonstrating the **end-to-end workflow** of a data scientist:

EDA → Feature Engineering → Modeling → Evaluation → Deployment.

We're excited to see your solutions, creativity, and technical skills in action. Good luck, and have fun with *The Test Case!* 